

# BUS445

2024-11-07

## Summary

Did some investigation for which variables are important. Removed NA missing values (4 total). Renamed some NULLS values to better explain why there were NULL. Did some feature engineering to look at continents, holiday, and seasonal results. I chose to remove the reservation\_status as the data is mainly after a cancellation/non cancellation has occurred.

Used random forests, decision tree, and logistic regression primarily. Random forests found that the top 5 important variables are: deposit\_type NonRefundable, country: Portugal, lead\_time, total\_of\_special\_requests, and previous\_cancellations. Decision trees found that the top 5 important variables are: deposit\_type Non Refundable, country Portugal, previous\_cancellation, market\_segment: Online TA and lead\_time. Quite similar due to random forests being the average of a bunch of decision trees.

Logistic regression was quite different and found company204, assigned\_room\_type1, agent252, company321, and required\_car\_parking\_spaces important. Probably due to logistic regression being linear function compared to the other methods.

Did 2 Neural networks since can't do one on all of the variables. One on the important variables from the random forest and decision tree. One on the important variables from logistic regression. The nn from the variables from the tree methods (RF and DT) performed pretty well. Important variables are deposit\_type Non Refund, previous\_cancellations, market\_segment Offline TA/TO, market\_segment Direct, and country PRT.

The nn on the variables from logistic regression performed pretty poorly. (Probably due to how it could only run with one hidden layer.) Important variables found are agent7, agent492, required\_car\_parking\_spaces, company48, agent28.

In terms of accuracy on the test set, Logistic Regression (84%) > Decision Tree (81%) > Neural Network Tree Variables (79%) > Random Forest (75%) > Neural Network Logistic Regression Variables (71%).

Still need to do plots on the important variables. Not sure if should do cross validation as it takes a while to run vs train test sets. Did a bit of train test sets at the end. Seems to perform better.

## Exploration of Missing Values

```
set.seed(123)
data=read.csv("hotel_bookings.csv")
originalData=data
#Checking for missing values (NA). Observed 4 missing values in the children column.
data[rowSums(is.na(data))>0,]
```

```

##          hotel is_canceled lead_time arrival_date_year arrival_date_month
## 40601 City Hotel          1          2          2015          August
## 40668 City Hotel          1          1          2015          August
## 40680 City Hotel          1          1          2015          August
## 41161 City Hotel          1          8          2015          August
##          arrival_date_week_number arrival_date_day_of_month
## 40601          32          3
## 40668          32          5
## 40680          32          5
## 41161          33          13
##          stays_in_weekend_nights stays_in_week_nights adults children babies meal
## 40601          1          0          2          NA          0          BB
## 40668          0          2          2          NA          0          BB
## 40680          0          2          3          NA          0          BB
## 41161          2          5          2          NA          0          BB
##          country market_segment distribution_channel is_repeated_guest
## 40601          PRT          Undefined          Undefined          0
## 40668          PRT          Direct          Undefined          0
## 40680          PRT          Undefined          Undefined          0
## 41161          PRT          Online TA          Undefined          0
##          previous_cancellations previous_bookings_not_canceled reserved_room_type
## 40601          0          0          B
## 40668          0          0          B
## 40680          0          0          B
## 41161          0          0          B
##          assigned_room_type booking_changes deposit_type agent company
## 40601          B          0          No Deposit          NULL          NULL
## 40668          B          0          No Deposit          14          NULL
## 40680          B          0          No Deposit          NULL          NULL
## 41161          B          0          No Deposit          9          NULL
##          days_in_waiting_list customer_type adr required_car_parking_spaces
## 40601          0 Transient-Party 12.0          0
## 40668          0 Transient-Party 12.0          0
## 40680          0 Transient-Party 18.0          0
## 41161          0 Transient-Party 76.5          0
##          total_of_special_requests reservation_status reservation_status_date
## 40601          1          Canceled          2015-08-01
## 40668          1          Canceled          2015-08-04
## 40680          2          Canceled          2015-08-04
## 41161          1          Canceled          2015-08-09

```

```

#Removing these 4 instances as there is a lot of observations
data=na.omit(data)

```

Contingency table of all the columns

```

#lapply(data,table) Commented out as it's too big of a print.

```

It's observed that there are NULL values in the data. The columns with NULL values are company, agent, and country.

```
colSums(data=="NULL")
```

```
##           hotel           is_canceled
##           0           0
##      lead_time      arrival_date_year
##           0           0
##      arrival_date_month      arrival_date_week_number
##           0           0
##      arrival_date_day_of_month      stays_in_weekend_nights
##           0           0
##      stays_in_week_nights      adults
##           0           0
##           children      babies
##           0           0
##           meal      country
##           0      488
##      market_segment      distribution_channel
##           0           0
##      is_repeated_guest      previous_cancellations
##           0           0
## previous_bookings_not_canceled      reserved_room_type
##           0           0
##      assigned_room_type      booking_changes
##           0           0
##           deposit_type      agent
##           0      16338
##           company      days_in_waiting_list
##      112589           0
##      customer_type      adr
##           0           0
##      required_car_parking_spaces      total_of_special_requests
##           0           0
##      reservation_status      reservation_status_date
##           0           0
```

The contingency table for the company feature.

```
#table(data$company) Commented out as it's too big of a print.
```

It is observed that the most common element is the NULL value with 112589 observations which is much more than 50% of the data. This is most likely due to a majority of the hotel bookings not be associated with a company booking. As a result, this implies that the NULL values are important so they will be renamed to "No Company"

```
data=data%>%mutate(company=ifelse(company=="NULL","No Company",company))
```

The agent feature has 16338 NULL values. As the agent number is related to the distribution channel of the booking, we will investigate the distribution channel.

```
#table(data$agent) Commented out as it's too big of a print.
```

```
agentNullData=data%>% filter(agent=="NULL")
#table(agentNullData$agent,agentNullData$distribution_channel) Commented out as it's too big of
a print.
```

Of the 16338 NULL values in the agent field, 13168 (5543+7625) of them belong to the corporate and direct distribution channels which have no agents as they directly contact the hotel for the booking. We will fill these with “No Travel Agency” as they don’t use any travel agency. There is 3167 NULL values with TA/TO distribution channels. We will fill in these with “TA/TO No Agent Number” as they have travel agents but have no agent id. The remaining 3 NULL values will be removed as they are only 3 of them.

```
data=data%>%mutate(agent=ifelse(distribution_channel %in% c("Corporate","Direct") & agent=='NUL
L','No Travel Agency',agent))
data=data%>%mutate(agent=ifelse(distribution_channel=="TA/TO" & agent=="NULL","TA/TO No Agent Nu
mber",agent))
data=data%>%filter(agent!="NULL")
```

Looking at the Contingency table of the country column we see that there is 488 NULL values.

```
#table(data$country) Commented out as it's too big of a print.
```

```
countryNulldata=data%>% filter(country=="NULL")
x=table(countryNulldata$country,countryNulldata$agent)
#x["NULL",] Commented out as it's too big of a print.
```

It is observed that majority of the observations with NULL for countries also had no agents which are now “No Travel Agency” and “TA/TO No Agent Number”. We will fill these with countries with “Unknown”. For all the other NULL countries, we will remove them as there is a small amount of them.

```
data=data%>%mutate(country=ifelse(agent %in% c("No Travel Agency","TA/TO No Agent Number") & cou
ntry=='NULL','Unknown',country))
data=data%>%filter(data$country!="NULL")
```

```
#lapply(data,table)
```

It is observed that there is 1168 undefined columns in the meal feature. As the other options are BB (Bed and Breakfast), FB(Full Board), HB(Half Board), and SC (Self Catering) it is observed that there is no option for no meal services. As a result, we will fill these undefined values with “Other”

```
data=data%>%mutate(meal=ifelse(meal=='Undefined','Other',meal))
table(data$meal)
```

```
##
##      BB      FB      HB Other      SC
## 92164    798 14450   1168 10649
```

```
head(data)
```

##	hotel	is_canceled	lead_time	arrival_date_year	arrival_date_month		
## 1	Resort Hotel	0	342	2015	July		
## 2	Resort Hotel	0	737	2015	July		
## 3	Resort Hotel	0	7	2015	July		
## 4	Resort Hotel	0	13	2015	July		
## 5	Resort Hotel	0	14	2015	July		
## 6	Resort Hotel	0	14	2015	July		
##	arrival_date_week_number	arrival_date_day_of_month	stays_in_weekend_nights				
## 1	27	1	0				
## 2	27	1	0				
## 3	27	1	0				
## 4	27	1	0				
## 5	27	1	0				
## 6	27	1	0				
##	stays_in_week_nights	adults	children	babies	meal	country	market_segment
## 1	0	2	0	0	BB	PRT	Direct
## 2	0	2	0	0	BB	PRT	Direct
## 3	1	1	0	0	BB	GBR	Direct
## 4	1	1	0	0	BB	GBR	Corporate
## 5	2	2	0	0	BB	GBR	Online TA
## 6	2	2	0	0	BB	GBR	Online TA
##	distribution_channel	is_repeated_guest	previous_cancellations				
## 1	Direct	0	0				
## 2	Direct	0	0				
## 3	Direct	0	0				
## 4	Corporate	0	0				
## 5	TA/TO	0	0				
## 6	TA/TO	0	0				
##	previous_bookings_not_canceled	reserved_room_type	assigned_room_type				
## 1	0	C	C				
## 2	0	C	C				
## 3	0	A	C				
## 4	0	A	A				
## 5	0	A	A				
## 6	0	A	A				
##	booking_changes	deposit_type	agent	company	days_in_waiting_list		
## 1	3	No Deposit	No Travel Agency	No Company	0		
## 2	4	No Deposit	No Travel Agency	No Company	0		
## 3	0	No Deposit	No Travel Agency	No Company	0		
## 4	0	No Deposit	304	No Company	0		
## 5	0	No Deposit	240	No Company	0		
## 6	0	No Deposit	240	No Company	0		
##	customer_type	adr	required_car_parking_spaces	total_of_special_requests			
## 1	Transient	0	0	0			
## 2	Transient	0	0	0			
## 3	Transient	75	0	0			
## 4	Transient	75	0	0			
## 5	Transient	98	0	1			
## 6	Transient	98	0	1			
##	reservation_status	reservation_status_date					
## 1	Check-Out	2015-07-01					
## 2	Check-Out	2015-07-01					

```
## 3      Check-Out      2015-07-02
## 4      Check-Out      2015-07-02
## 5      Check-Out      2015-07-03
## 6      Check-Out      2015-07-03
```


```
write.csv(data,"data.csv",row.names = FALSE) # Writing out for easier factor conversion
```

```
data=read.csv("data.csv",stringsAsFactors = TRUE)
data$is_canceled=as.factor(data$is_canceled)
file.remove("data.csv")
```

```
## [1] TRUE
```

## Correlation Exploration

```
library(corrplot)
numericData=data[sapply(data,is.numeric)]
corr=cor(numericData)
corrplot(corr,method="color")
```



As

there isn't any highly correlated columns, no columns will be removed.

# Creating New Features

Binning the lead time into quartiles

```
q1LeadTime=quantile(data$lead_time,0.25)
q2LeadTime=quantile(data$lead_time,0.50)
q3LeadTime=quantile(data$lead_time,0.75)
data$lead_timeCategories=cut(data$lead_time,breaks=c(-Inf,q1LeadTime,q2LeadTime,q3LeadTime,Inf),
labels=c("Very Low Lead Time", "Below Average Lead Time", "Above Average Lead Time", "High Lead Time"))
```

Making a continent Column

```
data$Continent=countrycode(data$country, origin = "iso3c", destination = "continent")
```

```
## Warning: Some values were not matched unambiguously: ATF, CN, TMP, UMI, Unknown
```

```
southAmerica=c("ARG", "BRA", "CHL", "PER", "COL", "VEN", "SUR", "ECU", "GUY", "PRY", "BOL", "GUY")
```

```
#Manually fixing continent values that the country code couldn't define
```

```
#South America is linked together as Americas with North America
```

```
data$Continent=ifelse(data$country %in% southAmerica & data$Continent == "Americas","South America",data$Continent)
```

```
data$Continent=ifelse(data$country == "ATF", "None",data$Continent) #French South Territories isn't associated with a continent
```

```
data$Continent=ifelse(data$country == "CN","Asia",data$Continent) #China
```

```
data$Continent=ifelse(data$country == "TMP","Asia",data$Continent) #East Timor, part of ASIA
```

```
data$Continent=ifelse(data$country == "UMI","None",data$Continent) #United States Minor Outlying Islands isn't associated with a continent
```

```
data$Continent=ifelse(data$country == "Unknown","Unknown",data$Continent)
```

Making a holiday seasons column (Summer, Christmas, New years)

```
data$ArrivalHolidaySeason=cut(data$arrival_date_week_number,breaks=c(-Inf,1,20,26,47,51,Inf),labels=c("New Year","Regular","Summer","Regular","Christmas","New Year"))
```

Making a seasonal column

```
data=data%>%mutate(ArrivalSeason=case_when(
  arrival_date_month %in% c("December", "January", "February") ~ "Winter",
  arrival_date_month %in% c("March", "April", "May") ~ "Spring",
  arrival_date_month %in% c("June", "July", "August") ~ "Summer",
  arrival_date_month %in% c("September", "October", "November") ~ "Fall")
)
data$ArrivalSeason=as.factor(data$ArrivalSeason)
```

```
originalData=data#Before removing columns stored original with features engineered for later use.
```

```
data=subset(data,select=-reservation_status) #Dropping variables that are observed after a hotel booking is finalized (Canceled, No Show, etc)
data=subset(data,select=-reservation_status_date)
```

```
data=subset(data,select=-arrival_date_week_number)#Dropping arrival week number as I used it to create the Seasonal columns
```

## Splitting the data for ML

```
data$is_canceled=as.factor(data$is_canceled)
data$Continent=as.factor(data$Continent)
partition=createDataPartition(data$is_canceled,p=0.75,list=FALSE)
data_train=data[partition,]
data_test=data[-partition,]
```

## Random Forest

Used cross validation to tune for parameters for RF.

```
#Commented out as it takes a while to run. The final values used for the model were mtry = 6, splitrule = gini and min.node.size = 10.
#gridRF=expand.grid(mtry=round(sqrt(ncol(data_train))),splitrule="gini",min.node.size= c(1, 5, 10, 20, 50))
#control=trainControl(method="cv",number=5,verboseIter=TRUE)
#model=train(is_canceled~.,data=data_train,method="ranger",tuneGrid=gridRF,trControl=control,importance = "impurity",num.trees=1000)
#varImp(model)
```

```
gridRF=expand.grid(mtry=6,splitrule="gini",min.node.size=10)
control=trainControl(method="cv",number=5)
rfmodel=train(is_canceled~.,data=data_train,method="ranger",tuneGrid=gridRF,trControl=control,importance = "impurity",num.trees=1000)
```

Accuracy

```
rf_preds=predict(rfmodel,newdata=data_test)
mean(rf_preds==data_test$is_canceled)
```

```
## [1] 0.7492871
```

Confusion Matrix

```
table(rf_preds,data_test$is_canceled)
```



```
##
## rf_preds      0      1
##           0 18748  7461
##           1   12  3586
```

## Variable Importance

```
rfImportance=varImp(rfmodel)
Top5RfImportance=rfImportance$importance%>%as.data.frame()%>%rownames_to_column("Feature") %>% arrange(desc(Overall))%>%head(5)
Top5RfImportance
```


```
##           Feature Overall
## 1 deposit_typeNon Refund 100.00000
## 2           countryPRT  53.65040
## 3           lead_time  33.83009
## 4 total_of_special_requests 30.54662
## 5 previous_cancellations 29.08125
```

*#Found Deposit type:Non refundable, country:Portugal, lead\_time, total of special requests, and previous\_cancellations and lead\_time important*

## Variable Importance Plot

```
ggplot(data=Top5RfImportance,mapping=aes(x=Overall,y= reorder(Feature, Overall)))+geom_bar(stat="identity",fill="blue")+scale_y_discrete(labels=c("Previous Cancellations","Total of Special Requests","Lead Time","Country: Portugal","Non Refundable Deposit"))+xlab("Importance")+ylab("Variables")+ggtitle("Most Important Variables from the Random Forest Model")
```

## Most Important Variables from the Random Forest Model



## Decision Tree

As with Random Forests, used cross validation to tune the parameters.

```
#Commented out as it takes a while to tune for the parameters. The final value used for the model was cp = 0.01  
#tune_gridTree=expand.grid(cp=seq(0.01,0.1, by=0.01))  
#train_controlTree=trainControl(method="cv",number=5,verboseIter=TRUE)  
#TreeModel=train(is_canceled~.,data=data_train,method="rpart",trControl=train_controlTree, tuneGrid = tune_gridTree)
```

```
tune_gridTree=expand.grid(cp=0.01)  
train_controlTree=trainControl(method="cv",number=5)  
TreeModel=train(is_canceled~.,data=data_train,method="rpart",trControl=train_controlTree, tuneGrid = tune_gridTree)  
Treepreds=predict(TreeModel,newdata=data_test)
```

Accuracy

```
mean(Treepreds==data_test$is_canceled)
```

```
## [1] 0.8131647
```

## Confusion Matrix

```
table(Treepreds,data_test$is_canceled)
```

```
##  
## Treepreds      0      1  
##           0 16210  3019  
##           1  2550  8028
```

```
TreeImportance=varImp(TreeModel)
```

## Important Variables

```
Top5TreeImportance=TreeImportance$importance%>%as.data.frame()%>%rownames_to_column("Feature") %  
>% arrange(desc(Overall))%>%head(5)  
Top5TreeImportance
```

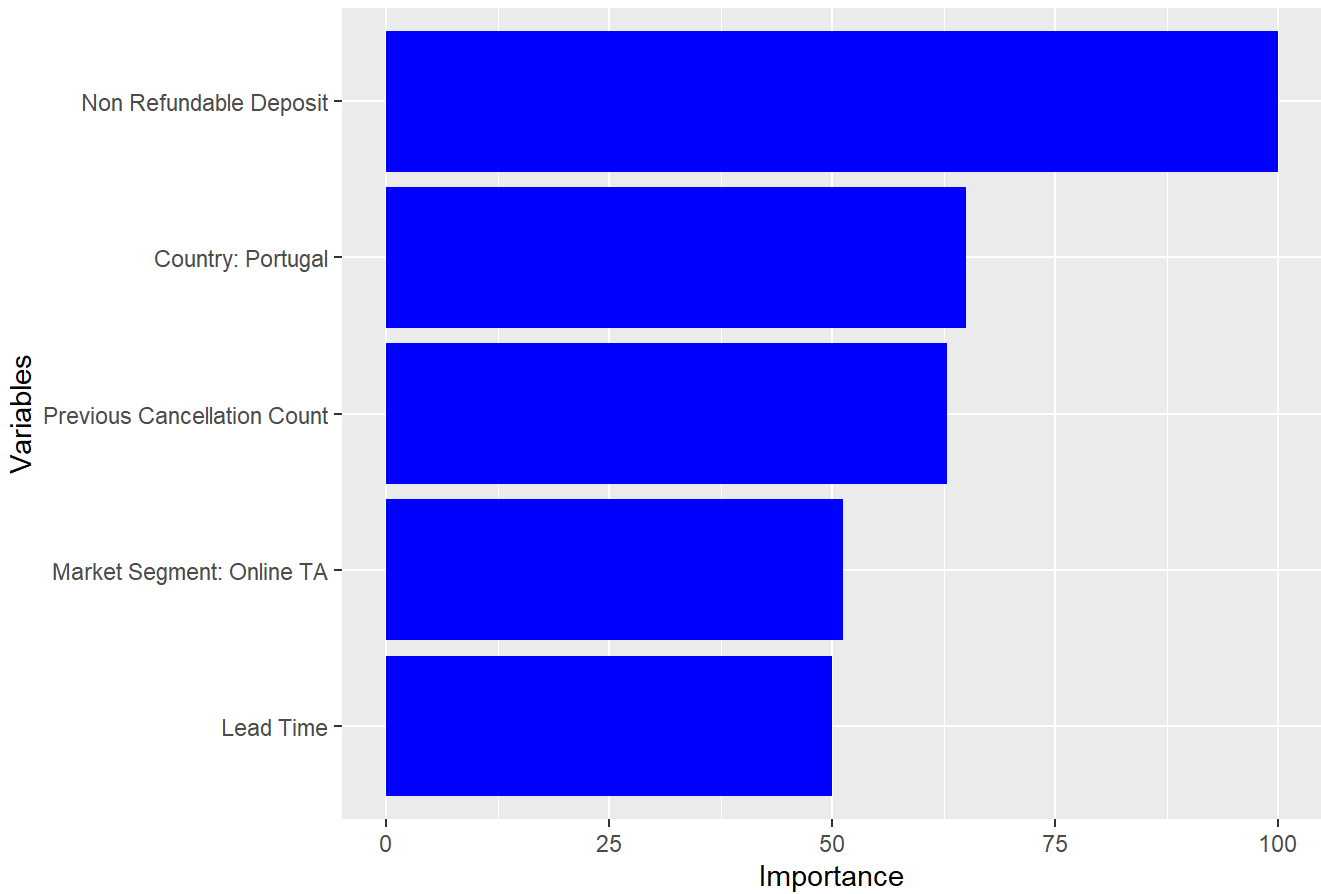
```
##           Feature Overall  
## 1 deposit_typeNon Refund 100.00000  
## 2           countryPRT  65.03180  
## 3 previous_cancellations 62.85566  
## 4 market_segmentOnline TA  51.15788  
## 5           lead_time  49.92142
```

```
#Found Deposity type:Non refundable, country:Portugal, previous_cancellations, Market Segment:On  
line TA, and Lead_time important
```

## Important Variables plot

```
ggplot(data=Top5TreeImportance,mapping=aes(x=Overall,y= reorder(Feature, Overall)))+geom_bar(sta  
t="identity",fill="blue")+scale_y_discrete(labels=c("Lead Time","Market Segment: Online TA","Pre  
vious Cancellation Count","Country: Portugal","Non Refundable Deposit"))+xlab("Importance")+ylab  
("Variables")+ggtitle("Most Important Variables from the Decision Tree Model")
```

Most Important Variables from the Decision Tree Model



## Logistic Regression

Was gonna use glm but it wouldn't run. So used multinom instead. Takes a while to run.

```
train_control=trainControl(method="cv",number=5)
LogitModel=train(is_canceled~.,data=data_train,method="multinom",trControl=train_control)
```

```
## # weights:  947 (946 variable)
## initial  value 49586.363003
## iter   10 value 35130.282435
## iter   20 value 32863.802475
## iter   30 value 31530.438436
## iter   40 value 29120.515376
## iter   50 value 27399.954803
## iter   60 value 26457.063805
## iter   70 value 25329.907761
## iter   80 value 24954.890807
## iter   90 value 24922.557808
## iter  100 value 24887.106597
## final   value 24887.106597
## stopped after 100 iterations
## # weights:  947 (946 variable)
## initial  value 49586.363003
## iter   10 value 35131.616109
## iter   20 value 32868.943902
## iter   30 value 31550.142755
## iter   40 value 29257.492742
## iter   50 value 27647.161521
## iter   60 value 26504.867007
## iter   70 value 25376.870077
## iter   80 value 25050.185836
## iter   90 value 25018.079964
## iter  100 value 24996.557525
## final   value 24996.557525
## stopped after 100 iterations
## # weights:  947 (946 variable)
## initial  value 49586.363003
## iter   10 value 35130.283767
## iter   20 value 32863.807579
## iter   30 value 31530.458143
## iter   40 value 29120.592000
## iter   50 value 27400.012965
## iter   60 value 26457.097132
## iter   70 value 25330.033852
## iter   80 value 24955.062251
## iter   90 value 24922.508283
## iter  100 value 24887.263721
## final   value 24887.263721
## stopped after 100 iterations
## # weights:  947 (946 variable)
## initial  value 49585.669856
## iter   10 value 33636.708238
## iter   20 value 31154.954915
## iter   30 value 30151.357492
## iter   40 value 28615.198547
## iter   50 value 26808.677740
## iter   60 value 26038.458354
## iter   70 value 25617.954283
## iter   80 value 25129.716465
```

```
## iter 90 value 24750.587557
## iter 100 value 24703.343942
## final value 24703.343942
## stopped after 100 iterations
## # weights: 947 (946 variable)
## initial value 49585.669856
## iter 10 value 33638.550142
## iter 20 value 31161.399562
## iter 30 value 30172.130823
## iter 40 value 28666.277212
## iter 50 value 27112.568121
## iter 60 value 26145.962930
## iter 70 value 25250.287108
## iter 80 value 25019.173422
## iter 90 value 24892.270172
## iter 100 value 24851.812200
## final value 24851.812200
## stopped after 100 iterations
## # weights: 947 (946 variable)
## initial value 49585.669856
## iter 10 value 33636.710110
## iter 20 value 31154.961730
## iter 30 value 30151.380721
## iter 40 value 28615.257688
## iter 50 value 26808.793815
## iter 60 value 26047.986347
## iter 70 value 25615.693912
## iter 80 value 25212.830809
## iter 90 value 24794.296059
## iter 100 value 24752.206068
## final value 24752.206068
## stopped after 100 iterations
## # weights: 947 (946 variable)
## initial value 49585.669856
## iter 10 value 34590.210632
## iter 20 value 32981.716428
## iter 30 value 32270.426818
## iter 40 value 30088.719669
## iter 50 value 28141.243414
## iter 60 value 27250.106544
## iter 70 value 25518.887691
## iter 80 value 25005.997912
## iter 90 value 24880.792430
## iter 100 value 24872.834126
## final value 24872.834126
## stopped after 100 iterations
## # weights: 947 (946 variable)
## initial value 49585.669856
## iter 10 value 34590.684472
## iter 20 value 32984.158575
## iter 30 value 32278.675305
## iter 40 value 30119.890093
```

```
## iter 50 value 28207.396786
## iter 60 value 27323.901460
## iter 70 value 25604.747913
## iter 80 value 25119.442918
## iter 90 value 24941.637885
## final value 24940.770390
## converged
## # weights: 947 (946 variable)
## initial value 49585.669856
## iter 10 value 34590.211105
## iter 20 value 32981.718850
## iter 30 value 32270.435001
## iter 40 value 30088.751934
## iter 50 value 28141.313103
## iter 60 value 27250.146714
## iter 70 value 25518.997685
## iter 80 value 25006.214417
## iter 90 value 24884.182214
## iter 100 value 24874.875718
## final value 24874.875718
## stopped after 100 iterations
## # weights: 947 (946 variable)
## initial value 49586.363003
## iter 10 value 33138.146033
## iter 20 value 30825.989087
## iter 30 value 29943.777784
## iter 40 value 28412.774451
## iter 50 value 27140.328427
## iter 60 value 26217.086123
## iter 70 value 25130.198695
## iter 80 value 24784.788052
## iter 90 value 24682.904002
## iter 100 value 24662.891933
## final value 24662.891933
## stopped after 100 iterations
## # weights: 947 (946 variable)
## initial value 49586.363003
## iter 10 value 33139.970792
## iter 20 value 30832.075871
## iter 30 value 29961.995101
## iter 40 value 28461.787162
## iter 50 value 27233.663495
## iter 60 value 26419.606143
## iter 70 value 25451.916770
## iter 80 value 24952.069108
## iter 90 value 24759.716875
## iter 100 value 24711.555745
## final value 24711.555745
## stopped after 100 iterations
## # weights: 947 (946 variable)
## initial value 49586.363003
## iter 10 value 33138.147822
```

```
## iter 20 value 30825.994753
## iter 30 value 29943.793773
## iter 40 value 28412.819375
## iter 50 value 27140.419642
## iter 60 value 26217.200665
## iter 70 value 25130.344284
## iter 80 value 24784.962877
## iter 90 value 24683.104299
## iter 100 value 24656.987505
## final value 24656.987505
## stopped after 100 iterations
## # weights: 947 (946 variable)
## initial value 49586.363003
## iter 10 value 34954.567017
## iter 20 value 32489.723584
## iter 30 value 31373.384046
## iter 40 value 29053.889967
## iter 50 value 27300.105836
## iter 60 value 26413.097206
## iter 70 value 25318.445725
## iter 80 value 24982.023299
## iter 90 value 24772.611531
## iter 100 value 24739.628827
## final value 24739.628827
## stopped after 100 iterations
## # weights: 947 (946 variable)
## initial value 49586.363003
## iter 10 value 34956.232471
## iter 20 value 32495.572822
## iter 30 value 31390.942122
## iter 40 value 29393.663034
## iter 50 value 27502.475240
## iter 60 value 26498.658652
## iter 70 value 25405.915796
## iter 80 value 25086.032423
## iter 90 value 24919.211107
## iter 100 value 24886.953251
## final value 24886.953251
## stopped after 100 iterations
## # weights: 947 (946 variable)
## initial value 49586.363003
## iter 10 value 34954.568681
## iter 20 value 32489.729416
## iter 30 value 31373.402137
## iter 40 value 29053.945801
## iter 50 value 27300.204452
## iter 60 value 26413.225819
## iter 70 value 25318.594343
## iter 80 value 24982.193450
## iter 90 value 24772.818413
## iter 100 value 24739.785320
## final value 24739.785320
```



```
## stopped after 100 iterations
## # weights:  947 (946 variable)
## initial  value 61982.607180
## iter   10 value 42306.598067
## iter   20 value 40200.162677
## iter   30 value 39495.687122
## iter   40 value 37263.275748
## iter   50 value 34531.822209
## iter   60 value 33469.930740
## iter   70 value 32297.303685
## iter   80 value 31356.902877
## iter   90 value 31058.200651
## iter  100 value 31055.281845
## final   value 31055.281845
## stopped after 100 iterations
```

```
Logitpreds=predict(LogitModel,newdata=data_test)
```

## Confusion Matrix

```
table(Logitpreds,data_test$is_canceled)
```

```
##
## Logitpreds      0      1
##           0 16891  2876
##           1  1869  8171
```

## Accuracy

```
mean(Logitpreds==data_test$is_canceled)
```

```
## [1] 0.8408092
```

## Important variables

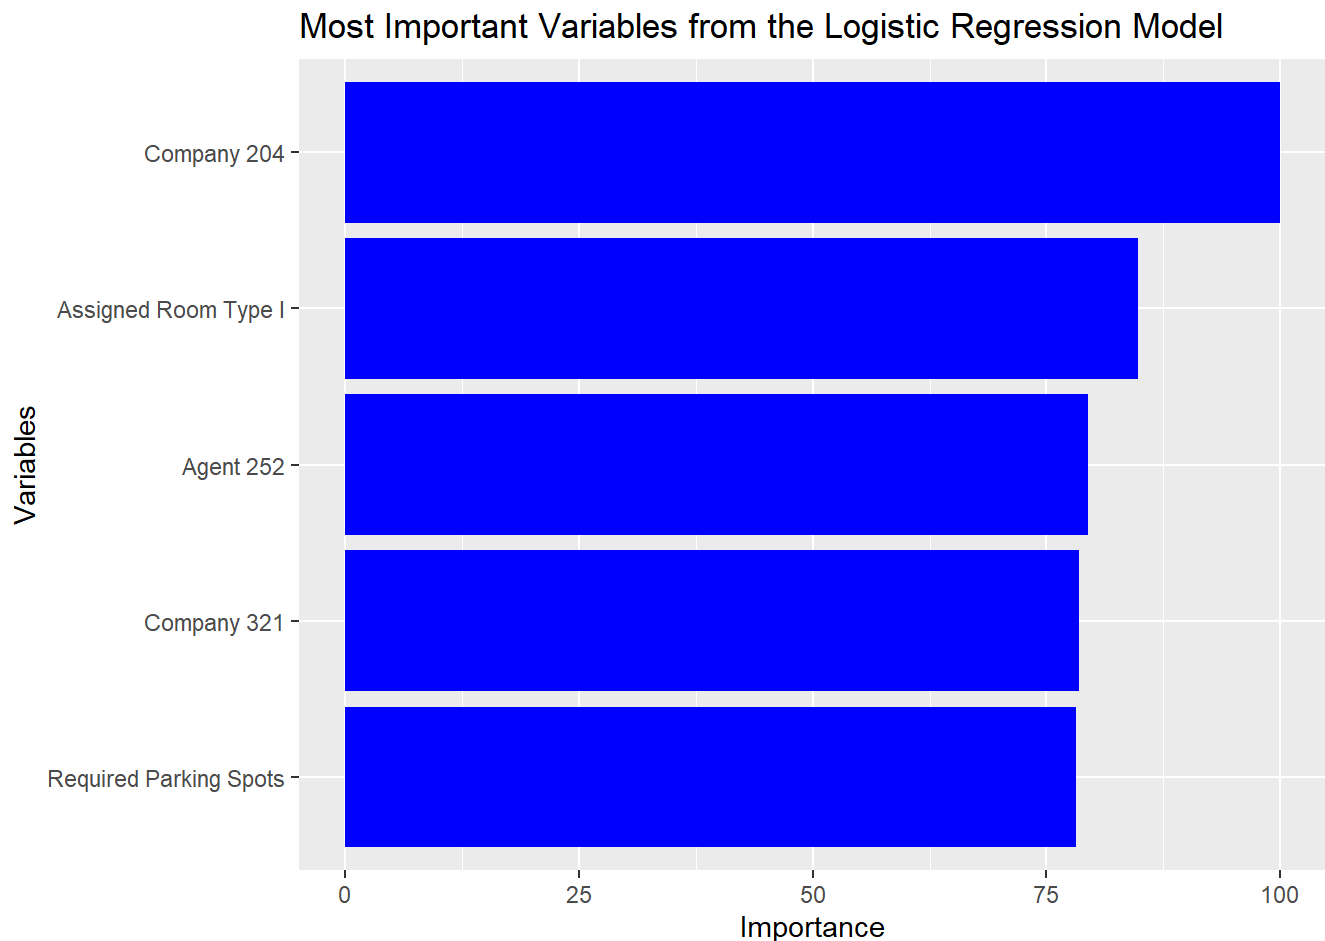
```
lgImportance=varImp(LogitModel)
Top5LgImportance=lgImportance$importance%>%as.data.frame()%>%rownames_to_column("Feature") %>% arrange(desc(Overall))%>%head(5)
Top5LgImportance
```

```
##      Feature Overall
## 1  agent252 100.00000
## 2 company321  84.83656
## 3  agent341  79.43836
## 4  agent276  78.52560
## 5 company204  78.15446
```

```
#Found company204, assigned_room_typeI, agent252, company321, and required_car_parking_spaces important
```

## Important Variables Plot

```
ggplot(data=Top5LgImportance,mapping=aes(x=Overall,y= reorder(Feature, Overall)))+geom_bar(stat="identity",fill="blue")+scale_y_discrete(labels=c("Required Parking Spots","Company 321","Agent 252","Assigned Room Type I","Company 204"))+xlab("Importance")+ylab("Variables")+ggtitle("Most Important Variables from the Logistic Regression Model")
```



Very different result from random forests and decision tree. Probably due to random forests struggling with linear relationships.

## Neural Net on variables from tree methods (RF and DT)

Can't use all the variables so used the ones that random forest and decision tree found important. Used CV to tune for parameters.

```
#size = 5 and decay = 0.1. Commented out as it takes too long to compute
#nn_trainControl=trainControl(method="cv",number=5,verboseIter = TRUE)
#nn_tuneGrid=expand.grid(size=c(1,2,3,5,10), decay = c(0.1, 0.2, 0.3))
#nnModel=train(is_canceled~lead_time+deposit_type+country+market_segment+previous_cancellations,
data=data_train,method="nnet",trControl=nn_trainControl,tuneGrid=nn_tuneGrid)

#only using the variables found as important in the previous sections as it gets too computation
ally complex
nn_trainControl=trainControl(method="cv",number=5)
nn_tuneGrid=expand.grid(size=5, decay = 0.1)
nnModel=train(is_canceled~lead_time+deposit_type+country+market_segment+previous_cancellations,d
ata=data_train,method="nnet",trControl=nn_trainControl,tuneGrid=nn_tuneGrid)
```

```
## # weights:  946
## initial  value 49453.927508
## iter   10 value 43863.460340
## iter   20 value 40382.798635
## iter   30 value 37505.261873
## iter   40 value 35236.559998
## iter   50 value 33608.339828
## iter   60 value 32861.702378
## iter   70 value 31819.216939
## iter   80 value 30934.205429
## iter   90 value 30428.354909
## iter  100 value 30033.687437
## final   value 30033.687437
## stopped after 100 iterations
## # weights:  946
## initial  value 48910.634061
## iter   10 value 44545.493973
## iter   20 value 37716.796804
## iter   30 value 34941.270849
## iter   40 value 31208.559049
## iter   50 value 30022.255571
## iter   60 value 29642.618032
## iter   70 value 29311.277651
## iter   80 value 29154.734835
## iter   90 value 29041.250737
## iter  100 value 28947.894981
## final   value 28947.894981
## stopped after 100 iterations
## # weights:  946
## initial  value 56286.614438
## iter   10 value 42205.296805
## iter   20 value 37985.206672
## iter   30 value 35697.429350
## iter   40 value 34460.406845
## iter   50 value 33065.625971
## iter   60 value 31455.782480
## iter   70 value 30957.721985
## iter   80 value 30520.923455
## iter   90 value 29855.698416
## iter  100 value 29549.725967
## final   value 29549.725967
## stopped after 100 iterations
## # weights:  946
## initial  value 48871.443807
## iter   10 value 43746.106173
## iter   20 value 42865.606983
## iter   30 value 39205.976855
## iter   40 value 34849.020670
## iter   50 value 33161.249913
## iter   60 value 32042.388785
## iter   70 value 31086.448700
## iter   80 value 30452.783919
```

```
## iter 90 value 30223.174408
## iter 100 value 29702.194831
## final value 29702.194831
## stopped after 100 iterations
## # weights: 946
## initial value 51121.708489
## iter 10 value 43479.016413
## iter 20 value 37609.681272
## iter 30 value 32645.483271
## iter 40 value 30269.527269
## iter 50 value 29687.528085
## iter 60 value 29305.195501
## iter 70 value 29090.062403
## iter 80 value 28963.168127
## iter 90 value 28826.093169
## iter 100 value 28749.071089
## final value 28749.071089
## stopped after 100 iterations
## # weights: 946
## initial value 71011.062020
## iter 10 value 54448.240675
## iter 20 value 51143.565351
## iter 30 value 42756.411149
## iter 40 value 40960.243892
## iter 50 value 38870.242087
## iter 60 value 38206.771824
## iter 70 value 37596.286203
## iter 80 value 37232.105544
## iter 90 value 36882.554511
## iter 100 value 36608.874204
## final value 36608.874204
## stopped after 100 iterations
```

```
nnPreds=predict(nnModel,newdata=data_test)
```

## Confusion Matrix

```
table(nnPreds,data_test$is_canceled)
```

```
##
## nnPreds      0      1
##      0 17361  4803
##      1  1399  6244
```

## Accuracy

```
mean(nnPreds==data_test$is_canceled)
```

```
## [1] 0.7919281
```

## Important variables

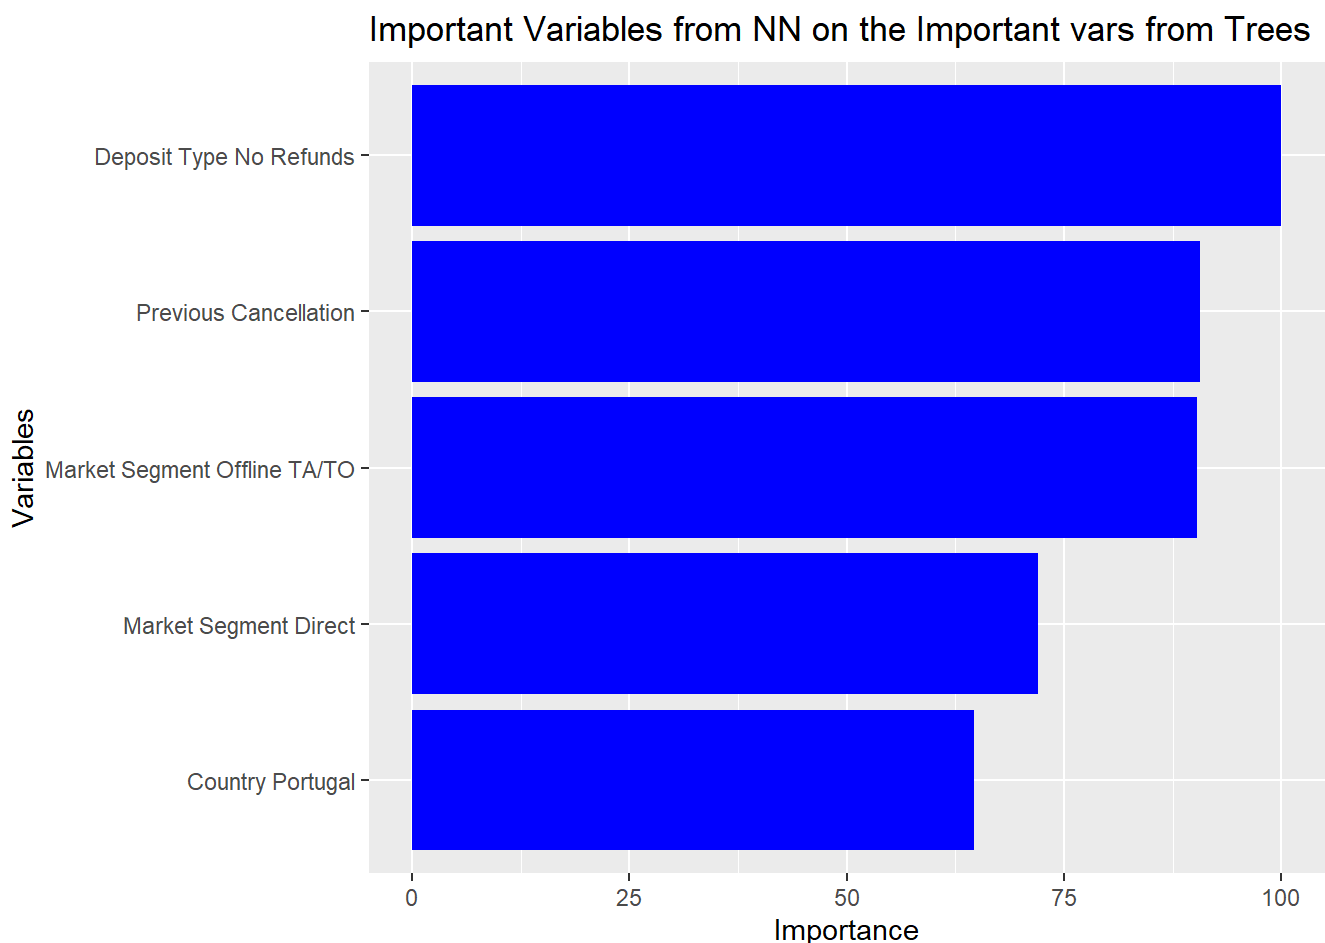
```
nnImportance=varImp(nnModel)
Top5nnImportance=nnImportance$importance%>%as.data.frame()%>%rownames_to_column("Feature") %>% arrange(desc(Overall))%>%head(5)
Top5nnImportance
```

```
##           Feature Overall
## 1 countryESP 100.00000
## 2 countryARE  90.65934
## 3 countryAGO  90.30891
## 4 market_segmentCorporate 72.05516
## 5 countryBRA  64.59718
```

```
#Found deposit_typeNon Refund, previous_cancellations, market_segmentOffline TA/TO, market_segmentDirect, countryPRT important
```

## Important Variables Plot

```
ggplot(data=Top5nnImportance,mapping=aes(x=Overall,y= reorder(Feature, Overall)))+geom_bar(stat="identity",fill="blue")+scale_y_discrete(labels=c("Country Portugal", "Market Segment Direct", "Market Segment Offline TA/TO", "Previous Cancellation", "Deposit Type No Refunds"))+xlab("Importance")+ylab("Variables")+ggtitle("Important Variables from NN on the Important vars from Trees")
```



# Neural Net on Important variables from Logistic Regression

```
#The final values used for the model were size = 1 and decay = 0.1. Commented out as it took too long to run
#NN_trainControl=trainControl(method="cv",number=5,verboseIter = TRUE)
#NN_tuneGrid=expand.grid(size=c(1,2), decay = c(0.1, 0.2, 0.3))
#NNModel=train(is_canceled~lead_time+deposit_type+country+market_segment+previous_cancellations,
data=data_train,method="nnet",trControl=NN_trainControl,tuneGrid=NN_tuneGrid, trace = FALSE)
#NNModel
```

```
NN_trainControl=trainControl(method="cv",number=5)
NN_tuneGrid=expand.grid(size=1, decay = 0.1)
NNModel=train(is_canceled~company+agent+assigned_room_type+required_car_parking_spaces,data=data_train,method="nnet",trControl=NN_trainControl,tuneGrid=NN_tuneGrid, trace = FALSE)
NNPreds=predict(NNModel,newdata=data_test)
```

## Confusion Matrix

```
table(NNPreds,data_test$is_canceled)
```

```
##
## NNPreds      0      1
##      0 16425  6199
##      1  2335  4848
```

## Accuracy

```
mean(NNPreds==data_test$is_canceled)
```

```
## [1] 0.7136914
```

## Important variables


```
NNImportance=varImp(NNModel)
Top5NNImportance=NNImportance$importance%>%as.data.frame()%>%rownames_to_column("Feature") %>% arrange(desc(Overall))%>%head(5)
Top5NNImportance
```

```
##           Feature Overall
## 1 assigned_room_typeK 100.00000
## 2 required_car_parking_spaces 75.22501
## 3 agent236 74.50452
## 4 agent220 63.59099
## 5 agent464 56.17624
```

```
#Found agent7, agent492, required_car_parking_spaces,company48,agent28 to be important
```

### Important Variables Plot

```
ggplot(data=Top5NNImportance,mapping=aes(x=Overall,y= reorder(Feature, Overall)))+geom_bar(stat="identity",fill="blue")+scale_y_discrete(labels=c("Agent 28","Company 48","Required Parking Spaces","Agent 492","Agent 7"))+xlab("Importance")+ylab("Variables")+ggtitle("Important Variables from the NN on Important Vars from LG")
```



Train test split Only. No cross validation.  
Computationally faster. Better Results too. Will explore more tmr.

```
rg=ranger(is_canceled~.,data=data_train,num.trees=1000,importance="impurity")
```

```
rgPreds=predict(rg,data=data_test)  
mean(rgPreds$predictions==data_test$is_canceled)
```

```
## [1] 0.8965679
```



```
importance(rg)
```

```
##          hotel          lead_time
##      393.04851      3841.73627
##      arrival_date_year      arrival_date_month
##      945.25155      1074.25613
##      arrival_date_day_of_month      stays_in_weekend_nights
##      1795.40185      748.70827
##      stays_in_week_nights      adults
##      1168.39678      434.71805
##      children      babies
##      229.69418      32.08636
##      meal      country
##      451.36078      4278.89217
##      market_segment      distribution_channel
##      1750.30362      395.76485
##      is_repeated_guest      previous_cancellations
##      73.87236      1232.20314
##      previous_bookings_not_canceled      reserved_room_type
##      143.92530      557.75932
##      assigned_room_type      booking_changes
##      923.09535      830.48841
##      deposit_type      agent
##      4565.07869      2234.72725
##      company      days_in_waiting_list
##      171.00354      75.43999
##      customer_type      adr
##      1008.58413      2552.53808
##      required_car_parking_spaces      total_of_special_requests
##      895.47107      2326.98966
##      lead_timeCategories      Continent
##      1488.44875      401.37790
##      ArrivalHolidaySeason      ArrivalSeason
##      386.83233      562.61466
```

```
table(rgPreds$predictions,data_test$is_canceled)
```

```
##
##      0      1
##  0 17595  1918
##  1  1165  9129
```

```
tree=rpart(is_canceled~.,data=data_train, method = "class")
```

```
treePreds=predict(tree,newdata=data_test,type="class")
mean(treePreds==data_test$is_canceled)
```

```
## [1] 0.8121582
```

```
table(treePreds,data_test$is_canceled)
```

```
##  
## treePreds      0      1  
##           0 17364  4203  
##           1  1396  6844
```

# Logistic Regression

```
lg=multinom(is_canceled~.,data=data_train,method="Binomial")
```

```
## # weights:  947 (946 variable)  
## initial  value 61982.607180  
## iter   10 value 42306.597569  
## iter   20 value 40200.160295  
## iter   30 value 39495.681300  
## iter   40 value 37263.253669  
## iter   50 value 34531.768919  
## iter   60 value 33469.863585  
## iter   70 value 32297.193898  
## iter   80 value 31356.737584  
## iter   90 value 31057.993696  
## final   value 31055.180283  
## converged
```

```
lgPreds=predict(lg,newdata=data_test)  
mean(lgPreds==data_test$is_canceled)
```

```
## [1] 0.8408092
```

```
table(lgPreds,data_test$is_canceled)
```

```
##  
## lgPreds      0      1  
##           0 16891  2876  
##           1  1869  8171
```